

Parallelsim: an accurate, generic, and efficient simulator for distributed deep learning

Regular Paper Published: 16 March 2026

Volume 8, pages 221–236 (2026) [Cite this article](#)



CCF Transactions on High
Performance Computing



Peng Liang, [Linbo Qiao](#), [Zhiquan Lai](#) & [Dongsheng Li](#) 

 121 Accesses [Explore all metrics](#) →

Abstract

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

[Manage preferences](#) for further information and to change your choices.

Accept all
cookies

Reject
optional
cookies

simulation. Finally, we evaluate ParallelSim on 16 DGX A100 nodes. The experimental results show that ParallelSim has an average prediction error of 1.83% on 16 nodes and helps to select the optimal strategies.

 This is a preview of subscription content, [log in via an institution](#)  to check access.

Access this article

[Log in via an institution](#) →

Subscribe and save

✓ Springer+

from ¥17,985 /Month

- Starting from 10 chapters or articles per month
- Access and download chapters and articles from more than 300k books and 2,500 journals
- Cancel anytime

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Explore related subjects

Discover the latest articles, books and news in related subjects, suggested using machine learning.

[Big Data](#)

[Computational platforms and environments](#)

[Numerical Simulation](#)

[Optimization](#)

[Register-Transfer-Level Implementation](#)

[Virtual Drug Screening](#)

[Deep Learning Frameworks and GPU Optimization](#)

References

Ansel, J, Yang, E, He, H, Gimelshein, N, Jain, A, Voznesensky, M., Bao, B, Bell, P, Berard, D, Burovski, E., Chauhan, G, Chourdia, A, Constable, W., Desmaison, A, DeVito, Z, Ellison, E, Feng, W., Gong, J, Gschwind, M, Hirsh, B., Huang, S, Kalambarkar, K, Kirsch, L, Lazos, M, Lezcano, M., Liang, Y, Liang, J, Lu, Y, Luk, C.K, Maher, B., Pan, Y, Puhersch, C, Reso, M, Saroufim, M, Siraichi, M.Y, Suk, H, Zhang, S, Suo, M, Tillet, P, Zhao, X, Wang, E, Zhou, K, Zou, R, Wang, X, Mathews, A, Wen, W, Chanan, G, Wu, P, Chintala, S: Pytorch 2: Faster machine learning through dynamic python bytecode transformation and graph compilation. In: Proceedings of the 29th ACM International Conference on Architectural

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Brown, T.B, Mann, B, Ryder, N, Subbiah, M, Kaplan, J, Dhariwal, P, Neelakantan, A, Shyam, P, Sastry, G, Askell, A, Agarwal, S, Herbert-Voss, A, Krueger, G, Henighan, T, Child, R, Ramesh, A, Ziegler, D.M, Wu, J, Winter, C, Hesse, C, Chen, M, Sigler, E, Litwin, M, Gray, S, Chess, B, Clark, J, Berner, C, McCandlish, S, Radford, A, Sutskever, I, Amodei, D: Language models are few-shot learners. In: Larochelle, H, Ranzato, M, Hadsell, R, Balcan, M, Lin, H. (eds.) Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, Virtual (2020). <https://doi.org/10.48550/arXiv.2005.14165>

Chen, Z, Deng, Y, Wu, Y, Gu, Q, Li, Y: Towards understanding the mixture-of-experts layer in deep learning. In: Koyejo, S, Mohamed, S, Agarwal, A, Belgrave, D, Cho, K, Oh, A. (eds.) Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022 (2022). http://papers.nips.cc/paper_files/paper/2022/hash/91edff07232fb1b55a505a9e9f6c0ff3-Abstract-Conference.html

Dao, T, Fu, D.Y, Ermon, S, Rudra, A, Ré, C: FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In: Advances in Neural Information Processing Systems

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [**privacy policy**](#) for more information on the use of your personal data.

for further information and to change your choices.

model. CoRR **abs/2405.04434** (2024)

<https://doi.org/10.48550/ARXIV.2405.04434>arXiv:2405.04434

Devlin, J, Chang, M, Lee, K, Toutanova, K.: BERT: pre-training of deep bidirectional transformers for language understanding. In: Burstein, J., Doran, C., Solorio, T. (eds.) Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2019, Minneapolis, MN, USA, June 2-7, 2019, Volume 1 (Long and Short Papers), pp. 4171–4186 (2019). <https://doi.org/10.18653/V1/N19-1423>

Duan, J., Li, X., Xu, P., Zhang, X., Yan, S., Liang, Y., Lin, D.: Proteus: Simulating the performance of distributed dnn training. IEEE Trans. Parallel Distrib. Syst. **35**(10), 1867–1878 (2024). <https://doi.org/10.1109/TPDS.2024.3443255>

Dubey, A, Jauhri, A, Pandey, A, Kadian, A, Al-Dahle, A, Letman, A, Mathur, A, Schelten, A, Yang, A, Fan, A, Goyal, A, Hartshorn, A, Yang, A, Mitra, A, Sravankumar, A, Korenev, A, Hinsvark, A, Rao, A, Zhang, A., Rodriguez, A, Gregerson, A, Spataru, A, Rozière, B, Biron, B, Tang, B, Chern, B, Caucheteux, C, Nayak, C, Bi, C, Marra, C, McConnell, C, Keller, C, Touret, C, Wu, C, Wong, C, Ferrer, C, Nikolaidis, C, Allonsius, D, Song, D, Pintz, D, Luvshits, D,

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Practice of Parallel Programming, Virtual Event, Republic of Korea, February 27- March 3, 2021, pp. 431–445 (2021). <https://doi.org/10.1145/3437801.3441593>

Fedus, W., Zoph, B., Shazeer, N.: Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. J. Mach. Learn. Res. **23**(120), 1–39 (2022)

He, J, Qiu, J, Zeng, A, Yang, Z, Zhai, J, Tang, J.: Fastmoe: A fast mixture-of-expert training system. CoRR **abs/2103.13262** (2021) <https://doi.org/10.48550/ARXIV.2103.13262>

He, H.: Strangely, Matrix Multiplications on GPUs Run Faster When Given "Predictable" Data. [Online]. Available: <https://www.thonking.ai/p/strangely-matrix-multiplications> (2024)

Huang, Y, Cheng, Y, Bapna, A, Firat, O, Chen, D, Chen, M.X, Lee, H, Ngiam, J, Le, Q.V, Wu, Y, Chen, Z.: Gpipe: Efficient training of giant neural networks using pipeline parallelism. In: Wallach, H.M., Larochelle, H., Beygelzimer, A, d'Alché-Buc, F, Fox, E.B, Garnett, R. (eds.) Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8–14, 2019, Vancouver, BC,

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Jacobs, S.A, Tanaka, M, Zhang, C, Zhang, M, Aminadabi, R.Y, Song, S.L, Rajbhandari, S, He, Y: System optimizations for enabling training of extreme long sequence transformer models. In: Proceedings of the 43rd ACM Symposium on Principles of Distributed Computing. PODC '24, pp. 121–130. Association for Computing Machinery, New York, NY, USA (2024). <https://doi.org/10.1145/3662158.3662806>

Korthikanti, V.A, Casper, J, Lym, S, McAfee, L, Andersch, M, Shoeybi, M, Catanzaro, B: Reducing activation recomputation in large transformer models (2023)

Lai, Z, Li, S, et al.: Merak: An efficient distributed dnn training framework with automated 3d parallelism for giant foundation models. IEEE Trans. Parallel Distrib. Syst., 1–13 (2023) <https://doi.org/10.1109/TPDS.2023.3247001>

Langer, M., He, Z., Rahayu, W., Xue, Y.: Distributed training of deep learning models: A taxonomic perspective. IEEE Trans. Parallel Distrib. Syst. **31**(12), 2802–2818 (2020). <https://doi.org/10.1109/TPDS.2020.3003307>

LeCun, Y., Bengio, Y., Hinton, G.E.: Deep learning. Nat. **521**(7553), 436–444 (2015).

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Li, S, Zhao, Y, Varma, R, Salpekar, O, Noordhuis, P, Li, T, Paszke, A, Smith, J, Vaughan, B, Damania, P, Chintala, S: Pytorch distributed: experiences on accelerating data parallel training. Proc. VLDB Endow. 13(12), 3005–3018 (2020)
<https://doi.org/10.14778/3415478.3415530>

Li, A., Song, S.L., Chen, J., Li, J., Liu, X., Tallent, N.R., Barker, K.J.: Evaluating modern gpu interconnect: Pcie, nvlink, nv-sli, nvswitch and gpudirect. IEEE Trans. Parallel Distrib. Syst. 31(1), 94–110 (2020). <https://doi.org/10.1109/TPDS.2019.2928289>

Liang, P., Tang, Y., Zhang, X., Bai, Y., Su, T., Lai, Z., Qiao, L., Li, D.: A survey on auto-parallelism of large-scale deep learning training. IEEE Trans. Parallel Distrib. Syst. 34(08), 2377–2390 (2023). <https://doi.org/10.1109/TPDS.2023.3281931>

Liu, Z, Cheng, S, Zhou, H, You, Y: Hanayo: Harnessing wave-like pipeline parallelism for enhanced large model training efficiency. The International Conference for High Performance Computing, Networking, Storage, and Analysis, 1–13 (2023)
<https://doi.org/10.1145/3581784.3607073>

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Lu, W, Yan, G, Li, J, Gong, S, Han, Y, Li, X: Flexflow: A flexible dataflow accelerator architecture for convolutional neural networks. In: 2017 IEEE International Symposium on High Performance Computer Architecture, HPCA 2017, Austin, TX, USA, February 4–8, 2017, pp. 553–564. IEEE Computer Society, ??? (2017).

<https://doi.org/10.1109/HPCA.2017.29>

Mayer, R., Jacobsen, H.: Scalable deep learning on distributed infrastructures: Challenges, techniques, and tools. ACM Comput. Surv. 53(1), 3–1337 (2021).

<https://doi.org/10.1145/3363554>

Narayanan, D, Phanishayee, A, Shi, K, Chen, X, Zaharia, M: Memory-efficient pipeline-parallel dnn training. In: International Conference on Machine Learning (ICML 2021) (2021). <https://doi.org/10.48550/arXiv.2006.09503>

Narayanan, D, Shoeybi, M, Casper, J, LeGresley, P, Patwary, M, Korthikanti, V, Vainbrand, D, Kashinkunti, P, Bernauer, J, Catanzaro, B, Phanishayee, A, Zaharia, M.: Efficient large-scale language model training on GPU clusters using megatron-lm. In: Supinski, B.R, Hall, M.W, Gamblin, T. (eds.) International Conference for High Performance Computing, Networking, Storage and Analysis, SC 2021, St. Louis, Missouri, USA, November 14–19, 2021, p. 58

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Patarasuk, P., Yuan, X.: Bandwidth optimal all-reduce algorithms for clusters of workstations. J. Parallel Distributed Comput. 69(2), 117–124 (2009).

<https://doi.org/10.1016/J.JPDC.2008.09.002>

Qi, H., Sparks, E.R., Talwalkar, A.: Paleo: A performance model for deep neural networks. In: 5th International Conference on Learning Representations, ICLR 2017, Toulon, France, April 24–26, 2017, Conference Track Proceedings (2017). <https://openreview.net/forum?id=SyVVJ85lg>

Qi, P., Wan, X., Huang, G., Lin, M.: Zero bubble (almost) pipeline parallelism. In: The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7–11, 2024 (2024). <https://openreview.net/forum?id=tuzTN0eIO5>

Rajbhandari, S., Rasley, J., Ruwase, O., He, Y.: Zero: memory optimizations toward training trillion parameter models. In: Cuicchi, C., Qualters, I., Kramer, W.T. (eds.) Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis, SC 2020, Virtual Event / Atlanta, Georgia, USA, November 9–19, 2020, p. 20. IEEE/ACM, ??? (2020). <https://doi.org/10.1109/SC41405.2020.00024>

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., Rodriguez, A., Joulin, A., Grave, E., Lample, G.: Llama: Open and efficient foundation language models. ArXiv [abs/2302.13971](https://arxiv.org/abs/2302.13971) (2023)
<https://doi.org/10.48550/arXiv.2302.13971>

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention is all you need. *Advances in neural information processing systems* **30** (2017)

Won, W., Heo, T., Rashidi, S., Sridharan, S., Srinivasan, S., Krishna, T.: Astra-sim2.0: Modeling hierarchical networks and disaggregated systems for large-model training at scale. In: 2023 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), pp. 283–294. IEEE Computer Society, Los Alamitos, CA, USA (2023).
<https://doi.org/10.1109/ISPASS57527.2023.00035>

Xiong, R., Yang, Y., He, D., Zheng, K., Zheng, S., Xing, C., Zhang, H., Lan, Y., Wang, L., Liu, T.: On layer normalization in the transformer architecture. In: Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13–18 July 2020, Virtual Event. *Proceedings of Machine Learning Research*, vol. 119, pp. 10524–10533 (2020)

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Processing Systems 2019, NeurIPS 2019, December 8–14, 2019, Vancouver, BC, Canada, pp. 12360–12371 (2019)

Zhao, Y., Gu, A., Varma, R., Luo, L., Huang, C.-C., Xu, M., Wright, L., Shojanazeri, H., Ott, M., Shleifer, S., Desmaison, A., Balioglu, C., Damania, P., Nguyen, B., Chauhan, G., Hao, Y., Mathews, A., Li, S.: Pytorch fsdp: Experiences on scaling fully sharded data parallel. Proc. VLDB Endow. 16(12), 3848–3860 (2023) <https://doi.org/10.14778/3611540.3611569>

Acknowledgements

We thank anonymous reviewers for their constructive comments. This work is sponsored in part by the National Natural Science Foundation of China under Grant No. 62025208 and 62421002.

Author information

Peng Liang, Linbo Qiao and Zhiquan Li have equally contributed to this work.

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Conflict of interest

On behalf of all authors, the corresponding author states that there is no Conflict of interest.

Rights and permissions

Springer Nature or its licensor (e.g. a society or other partner) holds exclusive rights to this article under a publishing agreement with the author(s) or other rightsholder(s); author self-archiving of the accepted manuscript version of this article is solely governed by the terms of such publishing agreement and applicable law.

[Reprints and permissions](#)

About this article

Cite this article

Liang, P., Qiao, L., Lai, Z. *et al.* Parallelsim: an accurate, generic, and efficient simulator for distributed deep learning. *CCF Trans. HPC* **8**, 221–236 (2026).

<https://doi.org/10.1007/s42514-025-00271-w>

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [privacy policy](#) for more information on the use of your personal data.

for further information and to change your choices.

Your privacy, your choice

We use essential cookies to make sure the site can function. We also use optional cookies for advertising, personalisation of content, usage analysis, and social media, as well as to allow video information to be shared for both marketing, analytics and editorial purposes.

By accepting optional cookies, you consent to the processing of your personal data - including transfers to third parties. Some third parties are outside of the European Economic Area, with varying standards of data protection.

See our [**privacy policy**](#) for more information on the use of your personal data.

for further information and to change your choices.